LoanLens

The thinking behind it

A thumbprint proves who signed. LoanLens checks that the borrower understood.

Artificial intelligence is the tool. Real intelligence is the brain.

The problem

A thumbprint proves who signed a loan. It does not prove the borrower understood it. The RBI's Key Fact Statement circular (RBI/2024-25/18) asks lenders to give the KFS in a language the borrower understands, explain its contents and obtain an acknowledgement of understanding. For someone who cannot read, that acknowledgement can still be a mark on a page they cannot read. LoanLens is built for that gap: it checks, out loud, whether they understood what is on the paper.

The idea

1. **Everything is said out loud**
   Every screen speaks its own instruction, so nobody has to read.

2. **Teach-back, not tick-box**
   The borrower explains the loan in their own words, with open questions in random order.

3. **One engine, three doors**
   In person, online with a KFS PDF or QR, and a call check as a stretch.

4. **A shared receipt**
   A PDF of what was checked, not a signature and not a loan approval.

How it works

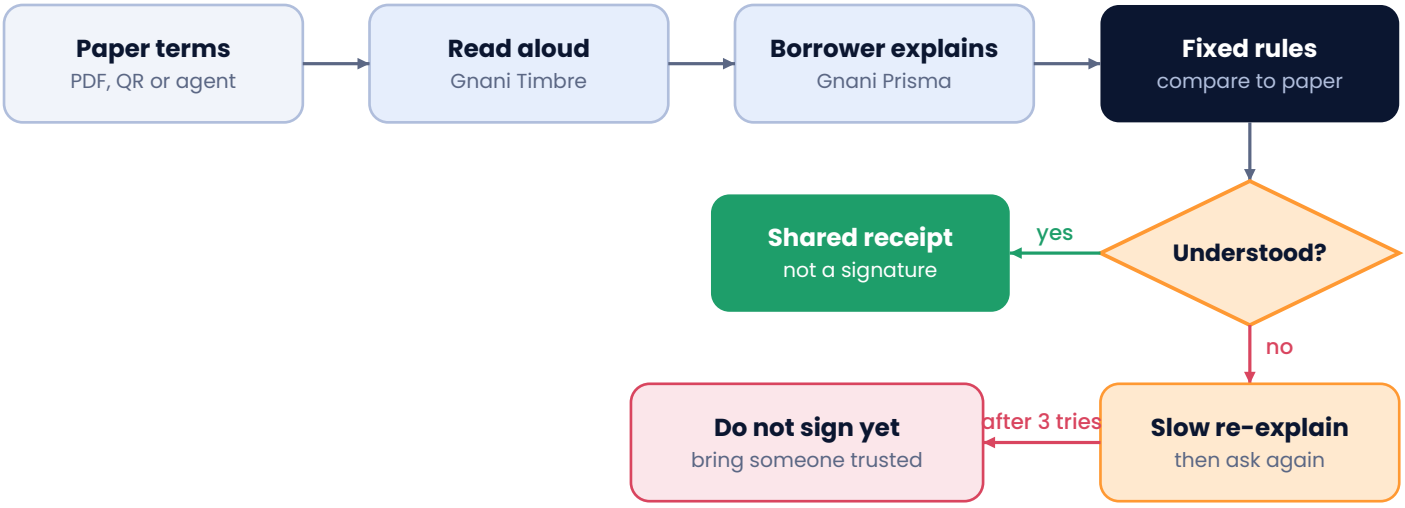

A passed gate gives proof of understanding. After 3 failed tries: do not sign yet. Fixed rules, not an AI model, decide.

Why we are different

Who is checked

Many tools explain terms or screen calls.
LoanLens checks whether the borrower understood, in their own words.

Who decides

Many tools let an AI model judge.
Fixed rules with unit tests decide. No AI judge, so every result traces to evidence.

How it is used

Many tools assume reading and typing.
Zero reading: every screen speaks, buttons are large, no typing.

Which models

Many tools mix several AI services.
Only Gnani: Prisma listens, Timbre speaks. Real calls, nothing faked.

What we researched

The rule behind the gap	The RBI Key Fact Statement circular (RBI/2024-25/18, 15 Apr 2024). It decided which terms we check: amount, rate, tenure, instalment, fees and late fee.
The Gnani APIs, read and tested	Prisma takes 30 seconds per clip and reads the format from the file name, so the browser encodes its own WAV. Prisma v3.0 returns numbers as digits. Only one request runs at a time, so all calls share one queue with retry. Timbre needs numbers written as words.
A real Hindi and English test	A spoken sample in the Gnani playground showed that the Hindi "nahi" comes after the field name. That exact output became a unit-test fixture.
The other entries	We looked at what others in the challenge built, to be sure we were solving a different problem: understanding, not explaining or scam-spotting.
Who we designed for	The zero-reading design, with sound plus colour plus icon for every status, comes from designing for someone who may not read the paper at all.

Consent and privacy, built in

Consent comes first.	In agent mode, the agent records a spoken consent line before anything else.
Recording is always visible.	A clear on-screen indicator shows when it records. Nothing records in the background.
Every clip is labelled.	Agent, borrower or helper, so it is clear who said what.
The borrower stays in control.	They can say "phir se boliye" to try again. There is no typing anywhere.
One button wipes it.	"Sab mitaayein" clears the whole session.
Private by design.	Synthetic data only. The API key stays on the server. Raw audio and document text are never logged.
Not a signature.	Voice is proof of understanding, alongside the signature.

Why it is built this way

Voice-first, zero reading	The borrower may not read. Every screen speaks, buttons are large, and status is colour plus icon plus sound.
Teach-back, not a quiz	Explaining in their own words shows more than tapping yes. Open questions come in random order.
No LLM decides anything	Money terms cannot be guessed. Fixed rules with 500+ unit tests make every comparison traceable to evidence.
Unknown stays unknown	A missing value is "not mentioned", never zero. A different unit or frequency is "couldn't check".
Only Gnani models	Prisma v3.0 listens and Timbre v2.5 speaks. Real calls only, with no other AI model in the app.
Careful words	Matched, different, not mentioned or couldn't check, plus the question to ask. Never "safe", "approved" or "fraud".

What we chose not to do

No AI judge · No voice biometrics · No voice cloning · No APR unless validated · No real data · No legal-signature claim

Honest limits	A prototype. Accuracy has not been measured on real borrowers. Hindi is the most tested language, Marathi wording is under review, and seven more languages are Beta and not enabled. All data is synthetic, and the call check is a stretch feature on synthetic calls only.